

1_generate_prediction

October 8, 2021

```
[1]: download_metadata = True # will create a list of all csv files in the s3 bucket  
      cache_http_calls = True # TTL for 1 hour
```

```
[ ]: # !pip install python-dotenv
```

```
[2]: import pyspark  
      import requests  
      import os  
      import pandas as pd  
      import boto3  
      import json  
      import cachetools  
  
      from botocore import UNSIGNED  
      from botocore.config import Config  
  
      from pyspark.sql.session import SparkSession  
      from pyspark.sql.types import *  
      from pyspark.sql.functions import *  
      from pyspark.sql import Row  
      from pyspark.conf import SparkConf  
  
      from copy import deepcopy  
      from datetime import datetime, timedelta  
      from dotenv import load_dotenv  
  
      load_dotenv()  
  
      os.environ['PYSPARK_SUBMIT_ARGS'] = '--packages "org.apache.hadoop:hadoop-aws:3.  
      ↪2.0" pyspark-shell'  
  
      from IPython.core.display import HTML  
      display(HTML("<style>pre { white-space: pre !important; }</style>"))  
  
      # !pip install boto3
```

```
# !pip install cachetools

sparkConf = SparkConf()
sparkConf.set("spark.hadoop.fs.s3a.aws.credentials.provider", "org.apache.
↳hadoop.fs.s3a.AnonymousAWSCredentialsProvider")
sparkConf.set("spark.hadoop.fs.s3a.threads.max", 10)
sparkConf.set("spark.hadoop.fs.s3a.endpoint", "s3.amazonaws.com")

sc = pyspark.SparkContext("local[*]", conf = sparkConf, appName = "
↳s_p_challenge")
spark = SparkSession(sc)

print(f"spark version = {spark.version}")
print(f"pyspark version = {pyspark.__version__}")
print(f"Hadoop version = {sc._jvm.org.apache.hadoop.util.VersionInfo.
↳getVersion()}")
```

<IPython.core.display.HTML object>

```
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.spark.unsafe.Platform
(file:/usr/local/spark-3.1.2-bin-hadoop3.2/jars/spark-unsafe_2.12-3.1.2.jar) to
constructor java.nio.DirectByteBuffer(long,int)
WARNING: Please consider reporting this to the maintainers of
org.apache.spark.unsafe.Platform
WARNING: Use --illegal-access=warn to enable warnings of further illegal
reflective access operations
WARNING: All illegal access operations will be denied in a future release

:: loading settings :: url = jar:file:/usr/local/spark-3.1.2-bin-
hadoop3.2/jars/ivy-2.4.0.jar!/org/apache/ivy/core/settings/ivysettings.xml

Ivy Default Cache set to: /home/jovyan/.ivy2/cache
The jars for the packages stored in: /home/jovyan/.ivy2/jars
org.apache.hadoop#hadoop-aws added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-
parent-564ea0b4-a118-4155-8ace-a65db8b3b76c;1.0
  confs: [default]
  found org.apache.hadoop#hadoop-aws;3.2.0 in central
  found com.amazonaws#aws-java-sdk-bundle;1.11.375 in central
:: resolution report :: resolve 169ms :: artifacts dl 2ms
  :: modules in use:
  com.amazonaws#aws-java-sdk-bundle;1.11.375 from central in [default]
  org.apache.hadoop#hadoop-aws;3.2.0 from central in [default]
```

	modules	artifacts
conf	number search dwnlded evicted	number dwnlded
default	2 0 0 0	2 0

```

-----
:: retrieving :: org.apache.spark#spark-submit-
parent-564ea0b4-a118-4155-8ace-a65db8b3b76c
  confs: [default]
    0 artifacts copied, 2 already retrieved (0kB/4ms)
21/10/08 14:30:25 WARN NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
setLogLevel(newLevel).

spark version = 3.1.2
pyspark version = 3.1.2
Hadoop version = 3.2.0

```

0.0.1 Download data and metadata from gdelt-open-data S3 bucket

```

[3]: # v2 event headers: https://github.com/linwoodc3/gdelt2HeaderRows/blob/master/
      ↪ schema_csvs/GDELT_2.0_Events_Column_Labels_Header_Row_Sep2016.csv
headers = pd.read_csv('headers.csv')
headers.head(n=2)

```

```

[3]:      tableId dataType      Empty \
0  GLOBALEVENTID  INTEGER  NULLABLE
1      SQLDATE  INTEGER  NULLABLE

                                     Description
0  Globally unique identifier assigned to each ev...
1  Date the event took place in YYYYMMDD format. ...

```

```

[5]: def download_metadata():
      """
      Download metadata from https://s3.console.aws.amazon.com/s3/buckets/
      ↪ gdelt-open-data?region=us-east-1
      """
      s3 = boto3.client('s3', config=Config(signature_version=UNSIGNED),
      ↪ region_name='us-east-1')
      s3_events = s3.list_objects_v2(Bucket='gdelt-open-data', Prefix='v2/events/
      ↪ ')
      s3_all_events = []

      is_truncated = True
      continuation_token = None

      while is_truncated:
          if continuation_token:

```

```

        s3_events = s3_events = s3.
        ↪list_objects_v2(Bucket='gdelt-open-data', Prefix='v2/events/',
        ↪ContinuationToken=continuation_token)
        else:
            s3_events = s3_events = s3.
            ↪list_objects_v2(Bucket='gdelt-open-data', Prefix='v2/events/')
            s3_all_events.append(s3_events)
            is_truncated = s3_events['IsTruncated']
            if 'NextContinuationToken' in s3_events:
                continuation_token = s3_events['NextContinuationToken']
            print('Total number of iterations to the S3 list objects = {:,}'.
            ↪format(len(s3_all_events)))
            s3_actual_events = []
            for s3_events in s3_all_events:
                s3_actual_events.extend(s3_events['Contents'])
            print('Total number of files in the S3 bucket = {:,}'.
            ↪format(len(s3_actual_events)))
            return s3_actual_events

```

```

[6]: if download_metadata:
    events_metadata = pd.DataFrame(download_metadata())
    events_metadata.sort_values(by='LastModified', inplace=True,
    ↪ascending=False)
    events_metadata

```

Total number of iterations to the S3 list objects = 144
 Total number of files in the S3 bucket = 143,462

```

[7]: events_metadata.head()

```

```

[7]:

```

	Key	LastModified	\
143461	v2/events/20190416151500.export.csv	2019-04-16 15:19:10+00:00	
143460	v2/events/20190416150000.export.csv	2019-04-16 15:03:11+00:00	
143459	v2/events/20190416144500.export.csv	2019-04-16 14:49:11+00:00	
143458	v2/events/20190416143000.export.csv	2019-04-16 14:34:10+00:00	
143457	v2/events/20190416141500.export.csv	2019-04-16 14:19:11+00:00	

	ETag	Size	StorageClass
143461	"eaec6933d33d8d4f437866a25f56c69d"	717467	STANDARD
143460	"4579b40b92ace49c845e5f042fcf12ec"	931359	STANDARD
143459	"10960e462cd359f563ce694f7e68d7c4"	738065	STANDARD
143458	"ded67702a20ef258d5a730343dd6918d"	794302	STANDARD
143457	"084ae0fab284cc18159e9d1dad2375ee"	713985	STANDARD

```

[8]: # Example uses GDELT dataset found here: https://aws.amazon.com/public-datasets/
    ↪gdelt/

```

```
events = spark.read.csv("s3a://gdelt-open-data/v2/events/20190416151500.export.
→csv", header=False, sep='\t', inferSchema=True)
print(f"Total number of events in current file: {events.count()}")
```

21/10/08 14:35:38 WARN MetricsConfig: Cannot locate configuration: tried hadoop-metrics2-s3a-file-system.properties,hadoop-metrics2.properties

Total number of events in current file: 1772

```
[9]: assert len(events.columns) == len(headers['tableId'])
for idx in range(len(events.columns)):
    events = events.withColumnRenamed(f"_c{idx}", list(headers['tableId'])[idx])
events = events.withColumn("SQLDATE", to_date(col("SQLDATE").cast("string"),
→"yyyyMMdd"))
events.printSchema()
events.show(n=2)
```

```
root
|-- GLOBALEVENTID: integer (nullable = true)
|-- SQLDATE: date (nullable = true)
|-- MonthYear: integer (nullable = true)
|-- Year: integer (nullable = true)
|-- FractionDate: double (nullable = true)
|-- Actor1Code: string (nullable = true)
|-- Actor1Name: string (nullable = true)
|-- Actor1CountryCode: string (nullable = true)
|-- Actor1KnownGroupCode: string (nullable = true)
|-- Actor1EthnicCode: string (nullable = true)
|-- Actor1Religion1Code: string (nullable = true)
|-- Actor1Religion2Code: string (nullable = true)
|-- Actor1Type1Code: string (nullable = true)
|-- Actor1Type2Code: string (nullable = true)
|-- Actor1Type3Code: string (nullable = true)
|-- Actor2Code: string (nullable = true)
|-- Actor2Name: string (nullable = true)
|-- Actor2CountryCode: string (nullable = true)
|-- Actor2KnownGroupCode: string (nullable = true)
|-- Actor2EthnicCode: string (nullable = true)
|-- Actor2Religion1Code: string (nullable = true)
|-- Actor2Religion2Code: string (nullable = true)
|-- Actor2Type1Code: string (nullable = true)
|-- Actor2Type2Code: string (nullable = true)
|-- Actor2Type3Code: string (nullable = true)
|-- IsRootEvent: integer (nullable = true)
|-- EventCode: integer (nullable = true)
|-- EventBaseCode: integer (nullable = true)
|-- EventRootCode: integer (nullable = true)
|-- QuadClass: integer (nullable = true)
```

```

|-- GoldsteinScale: double (nullable = true)
|-- NumMentions: integer (nullable = true)
|-- NumSources: integer (nullable = true)
|-- NumArticles: integer (nullable = true)
|-- AvgTone: double (nullable = true)
|-- Actor1Geo_Type: integer (nullable = true)
|-- Actor1Geo_FullName: string (nullable = true)
|-- Actor1Geo_CountryCode: string (nullable = true)
|-- Actor1Geo_ADM1Code: string (nullable = true)
|-- Actor1Geo_ADM2Code: string (nullable = true)
|-- Actor1Geo_Lat: double (nullable = true)
|-- Actor1Geo_Long: double (nullable = true)
|-- Actor1Geo_FeatureID: string (nullable = true)
|-- Actor2Geo_Type: integer (nullable = true)
|-- Actor2Geo_FullName: string (nullable = true)
|-- Actor2Geo_CountryCode: string (nullable = true)
|-- Actor2Geo_ADM1Code: string (nullable = true)
|-- Actor2Geo_ADM2Code: string (nullable = true)
|-- Actor2Geo_Lat: double (nullable = true)
|-- Actor2Geo_Long: double (nullable = true)
|-- Actor2Geo_FeatureID: string (nullable = true)
|-- ActionGeo_Type: integer (nullable = true)
|-- ActionGeo_FullName: string (nullable = true)
|-- ActionGeo_CountryCode: string (nullable = true)
|-- ActionGeo_ADM1Code: string (nullable = true)
|-- ActionGeo_ADM2Code: string (nullable = true)
|-- ActionGeo_Lat: double (nullable = true)
|-- ActionGeo_Long: double (nullable = true)
|-- ActionGeo_FeatureID: string (nullable = true)
|-- DATEADDED: long (nullable = true)
|-- SOURCEURL: string (nullable = true)

```

21/10/08 14:37:09 WARN package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

```

+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+

```

GLOBAL EVENT ID	SQL DATE	Month Year	Year	Fraction Date	Actor 1 Code	Actor 1 Name	Actor 1 Country Code	Actor 1 Known Group Code	Actor 1 Ethnic Code	Actor 1 Religion 1 Code	Actor 1 Religion 2 Code	Actor 1 Type 1 Code	Actor 1 Type 2 Code	Actor 1 Type 3 Code	Actor 2 Code	Actor 2 Name	Actor 2 Country Code	Actor 2 Known Group Code	Actor 2 Ethnic Code	Actor 2 Religion 1 Code	Actor 2 Religion 2 Code	Actor 2 Type 1 Code	Actor 2 Type 2 Code	Actor 2 Type 3 Code	Is Root Event	Event Code	Event Base Code	Event Root Code	Quad Class	Goldstein Scale	Num Mentions	Num Sources	Num Articles	Avg Tone	Actor 1 Geo_Type	Actor 1 Geo_FullName	Actor 1 Geo_CountryCode	Actor 1 Geo_ADM1Code	Actor 1 Geo_ADM2Code	Actor 1 Geo_Lat	Actor 1 Geo_Long	Actor 1 Geo_FeatureID	Actor 2 Geo_Type	Actor 2 Geo_FullName	Actor 2 Geo_CountryCode	Actor 2 Geo_ADM1Code	Actor 2 Geo_ADM2Code	Actor 2 Geo_Lat	Actor 2 Geo_Long	Actor 2 Geo_FeatureID	Action Geo_Type	Action Geo_FullName	Action Geo_CountryCode	Action Geo_ADM1Code	Action Geo_ADM2Code	Action Geo_Lat	Action Geo_Long	Action Geo_FeatureID	DATE ADDED	SOURCE URL
-----------------	----------	------------	------	---------------	--------------	--------------	----------------------	--------------------------	---------------------	-------------------------	-------------------------	---------------------	---------------------	---------------------	--------------	--------------	----------------------	--------------------------	---------------------	-------------------------	-------------------------	---------------------	---------------------	---------------------	---------------	------------	-----------------	-----------------	------------	-----------------	--------------	-------------	--------------	----------	------------------	----------------------	-------------------------	----------------------	----------------------	-----------------	------------------	-----------------------	------------------	----------------------	-------------------------	----------------------	----------------------	-----------------	------------------	-----------------------	-----------------	---------------------	------------------------	---------------------	---------------------	----------------	-----------------	----------------------	------------	------------

	838788879 2018-04-16	201804 2018	2018.2904	null	null
null	null	null	null	null	
null	null	null	null	IRL	IRELAND
IRL	null	null		null	
null	null	null	null	1	43
43	4	1	2.8	10	1
10 3.58208955223881		0	null		null
null	null	null	null		null
0	null	null		null	
null	null	null	null		0
null	null	null	null	null	null
null	20190416151500 https://www.eveni...				
	838788880 2018-04-16	201804 2018	2018.2904	BUS	INVESTOR
null	null	null		null	
null	BUS	null	null	null	null
null	null	null		null	
null	null	null	null	1	874
87	8	2	10.0	10	1
10 2.48520710059172		1	Germany		GM
GM	null	51.5	10.5		GM
0	null	null		null	

[illegible]

```
[10]: print(f"Total number of events before cleaning: {events.count()}")
events_clean = events.filter('Actor1Code is not Null and Actor2Code is not Null
↳ and Actor1Geo_Lat is not Null and Actor1Geo_Long is not Null and
↳ Actor2Geo_Lat is not Null and Actor2Geo_Long is not Null')
print(f"Total number of events after cleaning: {events_clean.count()}")
events_clean.show(n=2)
```

[illegible]


```

response = json.loads(res.content.decode('utf-8'))
call_cache[event_id] = response
return response

def flatten_model_response(actor: str, response: dict[str, object],
    ↪event_id=None, debug=False):
    d = {}
    if event_id and debug:
        print(event_id)
    try:
        d[f'{actor}_model_time_in_ms'] = response['model_time_in_ms']
        d[f'{actor}_release_harness_version'] = ↪
    ↪response['release']['harness_version']
        d[f'{actor}_release_model_version'] = ↪
    ↪response['release']['model_version']
        d[f'{actor}_release_model_version_number'] = ↪
    ↪response['release']['model_version_number']
        d[f'{actor}_request_id'] = response['request_id']
        d[f'{actor}_result_class1'] = response['result']['class1']
        d[f'{actor}_result_class2'] = response['result']['class2']
        d[f'{actor}_timing'] = response['timing']
    except Exception as e:
        print(response)
    return d

```

```

[12]: def create_payload(avg_tone, goldstein, actor_code, lat, lon, date):
    data = {}
    data['avg_tone'] = avg_tone
    data['goldstein'] = goldstein
    data['actor_code'] = actor_code
    data['lat'] = lat
    data['lon'] = lon
    data['date'] = date.strftime('%Y-%m-%d %H:%M:%S')
    payload = {}
    payload['data'] = data
    return payload

def call_model_output(row):
    '''
        payload = {
            "data": {
                "avg_tone": -2,
                "goldstein": 0.5,
                "actor_code": "GOV",
                "lat": 38,
                "lon": -78,

```

```

        "date": "2018-10-23 04:30:00"
    }
}

'''
# actor 1
r = row.asDict(True)
payload = create_payload(row['AvgTone'], row['GoldsteinScale'],
    row['Actor1Code'], row['Actor1Geo_Lat'], row['Actor1Geo_Long'], datetime.
    strptime(str(row['DATEADDED']), '%Y%m%d%H%M%S'))
response = flatten_model_response('Actor1', get_model_response(payload),
    event_id=row['GLOBALEVENTID'])

for k, v in response.items():
    r[k] = v

# actor 2
payload = create_payload(row['AvgTone'], row['GoldsteinScale'],
    row['Actor2Code'], row['Actor2Geo_Lat'], row['Actor2Geo_Long'], datetime.
    strptime(str(row['DATEADDED']), '%Y%m%d%H%M%S'))
response = flatten_model_response('Actor2', get_model_response(payload),
    event_id=row['GLOBALEVENTID'])

for k, v in response.items():
    r[k] = v

return Row(**r)

def define_schema(events):
    schema = deepcopy(events.schema)
    print('Number of columns in schema before addition = {:,}'.
    format(len(schema)))
    # https://spark.apache.org/docs/latest/sql-ref-datatypes.html
    for actor in ['Actor1', 'Actor2']:
        schema.add(StructField(f'{actor}__model_time_in_ms', IntegerType(),
    True))
        schema.add(StructField(f'{actor}_release_harness_version',
    StringType(), True))
        schema.add(StructField(f'{actor}_release_model_version', StringType(),
    True))
        schema.add(StructField(f'{actor}_release_model_version_number',
    IntegerType(), True))
        schema.add(StructField(f'{actor}_request_id', StringType(), True))
        schema.add(StructField(f'{actor}_result_class1', BooleanType(), True))
        schema.add(StructField(f'{actor}_result_class2', IntegerType(), True))
        schema.add(StructField(f'{actor}_timing', DoubleType(), True))

```

```

    print('Number of columns in schema after addition = {:,}'.
    ↪format(len(schema)))
    return schema

```

```

[13]: df = events_clean.rdd.map(call_model_output)
      schema = define_schema(events_clean)
      df = spark.createDataFrame(df, schema)
      df.show(n=1, vertical=True)
      df.write.parquet('model_output.parquet')

```

Number of columns in schema before addition = 61

Number of columns in schema after addition = 77

[Stage 10:>

(0 + 1) / 1]

```

-RECORD 0-----
GLOBALEVENTID          | 838788881
SQLDATE                 | 2018-04-16
MonthYear               | 201804
Year                   | 2018
FractionDate            | 2018.2904
Actor1Code              | EDU
Actor1Name              | ECONOMIST
Actor1CountryCode       | null
Actor1KnownGroupCode    | null
Actor1EthnicCode        | null
Actor1Religion1Code     | null
Actor1Religion2Code     | null
Actor1Type1Code         | EDU
Actor1Type2Code         | null
Actor1Type3Code         | null
Actor2Code              | GOV
Actor2Name              | REGULATOR
Actor2CountryCode       | null
Actor2KnownGroupCode    | null
Actor2EthnicCode        | null
Actor2Religion1Code     | null
Actor2Religion2Code     | null
Actor2Type1Code         | GOV
Actor2Type2Code         | null
Actor2Type3Code         | null
IsRootEvent             | 1
EventCode               | 20
EventBaseCode           | 20
EventRootCode           | 2
QuadClass               | 1
GoldsteinScale          | 3.0
NumMentions             | 10

```

NumSources	1
NumArticles	10
AvgTone	-3.15315315315315
Actor1Geo_Type	4
Actor1Geo_FullName	Vancouver, Britis...
Actor1Geo_CountryCode	CA
Actor1Geo_ADM1Code	CA02
Actor1Geo_ADM2Code	12552
Actor1Geo_Lat	49.25
Actor1Geo_Long	-123.133
Actor1Geo_FeatureID	-575268
Actor2Geo_Type	4
Actor2Geo_FullName	Vancouver, Britis...
Actor2Geo_CountryCode	CA
Actor2Geo_ADM1Code	CA02
Actor2Geo_ADM2Code	12552
Actor2Geo_Lat	49.25
Actor2Geo_Long	-123.133
Actor2Geo_FeatureID	-575268
ActionGeo_Type	4
ActionGeo_FullName	Vancouver, Britis...
ActionGeo_CountryCode	CA
ActionGeo_ADM1Code	CA02
ActionGeo_ADM2Code	12552
ActionGeo_Lat	49.25
ActionGeo_Long	-123.133
ActionGeo_FeatureID	-575268
DATEADDED	20190416151500
SOURCEURL	https://www.bnnbl...
Actor1__model_time_in_ms	0
Actor1_release_harness_version	0.1
Actor1_release_model_version	5ec427ae4cedfd000...
Actor1_release_model_version_number	4
Actor1_request_id	57W806FRXSHYBI3G
Actor1_result_class1	true
Actor1_result_class2	3
Actor1_timing	0.08416175842285156
Actor2__model_time_in_ms	0
Actor2_release_harness_version	0.1
Actor2_release_model_version	5ec427ae4cedfd000...
Actor2_release_model_version_number	4
Actor2_request_id	V7UYFM2WQTZ8SXHE
Actor2_result_class1	true
Actor2_result_class2	3
Actor2_timing	0.06723403930664062

only showing top 1 row

```

-----
AnalysisException                                Traceback (most recent call last)
/tmp/ipykernel_145/2354079784.py in <module>
      3 df = spark.createDataFrame(df, schema)
      4 df.show(n=1, vertical=True)
----> 5 df.write.parquet('model_output.parquet')

/usr/local/spark/python/pyspark/sql/readwriter.py in parquet(self, path, mode,
↳ partitionBy, compression)
    1248         self.partitionBy(partitionBy)
    1249         self._set_opts(compression=compression)
-> 1250         self._jwrite.parquet(path)
    1251
    1252     def text(self, path, compression=None, lineSep=None):

/usr/local/spark/python/lib/py4j-0.10.9-src.zip/py4j/java_gateway.py in
↳ __call__(self, *args)
    1302
    1303         answer = self.gateway_client.send_command(command)
-> 1304         return_value = get_return_value(
    1305             answer, self.gateway_client, self.target_id, self.name)
    1306

/usr/local/spark/python/pyspark/sql/utils.py in deco(*a, **kw)
    115         # Hide where the exception came from that shows a
↳ non-Pythonic
    116         # JVM exception message.
--> 117         raise converted from None
    118     else:
    119         raise

AnalysisException: path file:/home/jovyan/work/model_output.parquet already
↳ exists.

```

```

[ ]: df = spark.read.parquet('model_output.parquet')
     dfp = df.toPandas()
     dfp.to_csv('model_output.csv')

```